

Fall 2023 CS543/ECE549

Assignment 5: Affine factorization and binocular stereo

Due date: Wednesday, December 6, 11:59:59PM

- [Part 1: Affine factorization](#)
- [Part 2: Binocular stereo](#)

Part 1: Affine Factorization

The goal of this part of the assignment is to implement the Tomasi and Kanade affine structure from motion method as described in lecture. You will be working with Carlo Tomasi's 101-frame hotel sequence:



[Download](#) the data file, including all the 101 images and a measurement matrix consisting of 215 points visible in each of the 101 frames (see readme file inside archive for details).

1. Load the data matrix and normalize the point coordinates by translating them to the mean of the points in each view (see lecture for details).
2. Apply SVD to the $2M \times N$ data matrix to express it as $D = U @ W @ V'$ (using NumPy notation) where U is a $2M \times 3$ matrix, W is a 3×3 matrix of the top three singular values, and V is a $N \times 3$ matrix. You can use `numpy.linalg.svd` to compute this decomposition. Next, derive structure and motion matrices from the SVD as explained in the lecture.
3. Find the matrix Q to eliminate the affine ambiguity using the method described on slide 32 of the lecture.
4. Use `matplotlib` to display the 3D structure (in your report, you may want to include snapshots from several viewpoints to show the structure clearly). Discuss whether or not the reconstruction has an ambiguity.
5. Display three frames with both the observed feature points and the estimated projected 3D points overlaid. Report your total residual (sum of squared Euclidean distances, in pixels, between the observed and the reprojected features) over all the frames, and plot the per-frame residual as a function of the frame number.

Part 1 Extra Credit

- Incorporate incomplete point tracks (i.e., points that are not visible in every frame) into the reconstruction. Use the tracks from [this file](#).
- Create a textured 3D model of the reconstructed points. For example, you can compute a Delaunay triangulation of the points in 2D, then lift that mesh structure to 3D, and texture it using the images.

Part 2: Binocular Stereo

The goal of this part is to implement a simple window-based stereo matching algorithm for rectified stereo pairs. You will be using the following stereo pairs:



Follow the basic outline given in lecture: pick a window around each pixel in the first (reference) image, and then search the corresponding scanline in the second image for a matching window. The output should be a disparity map with respect to the first view (use these ground truth maps for qualitative reference for [first pair](#) and [second pair](#)). You should experiment with the following settings and parameters:

- **Search window size:** show disparity maps for several window sizes and discuss which window size works the best (or what are the tradeoffs between using different window sizes). How does the running time depend on window size?
- **Disparity range:** what is the range of the scanline in the second image that should be traversed in order to find a match for a given location in the first image? Examine the stereo pair to determine what is the maximum disparity value that makes sense, where to start the search on the scanline, and which direction to search in. Report which settings you ended up using.
- **Matching function:** try sum of squared differences (SSD), sum of absolute differences (SAD), and normalized correlation. Discuss in your report whether there is any difference between using these functions, both in terms of quality of the results and in terms of running time.

In addition to showing your results and discussing implementation parameters, discuss the shortcomings of your algorithm. Where do the estimated disparity maps look good, and where do they look bad? What would be required to produce better results? Also discuss the running time of your approach and what might be needed to make stereo run faster.

Part 2 Extra Credit

- Convert your disparity map to a depth map and attempt to **visualize the depth map in 3D**. Just pretend that all projection rays are parallel, and try to "guess" the depth scaling constant. Experiment with displaying a 3D point cloud, or computing a Delaunay triangulation of that point cloud.
- Find **additional rectified stereo pairs** on the Web and show the results of your algorithm on these pairs.
- Find **non-rectified stereo pairs and rectification code** on the Web and apply your algorithm to this data.
- Implement **multiple-baseline stereo** as described in [this paper](#). Use [this data](#).
- Try to incorporate **non-local constraints** (smoothness, uniqueness, ordering) into your algorithm. You can come up with simple heuristic ways of incorporating these constraints, or try to implement some of the more advanced methods discussed in the course (dynamic programming, graph cuts). For this part, it is also fine to find code on the web.

Submission Instructions

You must upload the following files on [Canvas](#):

1. Your code in two separate files for part 1 and part 2. The filenames should be **lastname_firstname_a5_p1.py** and **lastname_firstname_a5_p2.py**. We prefer that you upload .py python files, but if you use a Python notebook, make sure you upload both the original .ipynb file and an exported PDF of the notebook.
2. A report **in a single PDF file** with all your results and discussion for both parts following this [template](#). The filename should be **lastname_firstname_a5.pdf**.
3. All your output images and visualizations **in a single zip file**. The filename should be **lastname_firstname_a5.zip**. Note that this zip file is for backup documentation only, in case we cannot see the images in your PDF report clearly enough. **You will not receive credit for any output images that are part of the zip file but are not shown (in some form) in the report PDF.**

Please refer to [course policies](#) on academic honesty, collaboration, late days, etc.